

Regression analysis I

Instructor: Li, Han

Contents

- ❑ The framework of OLS regression
- ❑ Fit the data using linear model
- ❑ Fit the data using quadratic model
- ❑ Analyze the output results
- ❑ Check the model assumptions

Regression analysis

Regression analysis lives at the heart of statistics. It is a broad term for a set of methodologies used to analyze the functional relationship between a response variable and one or more explanatory(predictor) variables.

In general, regression analysis can

- ▣ identify the relevant explanatory variables
- ▣ describe the form of the relationships involved
- ▣ predict the response for given explanatory variables.

A data example

An exercise physiologist might use regression analysis to develop an equation for predicting the expected number of calories a person will burn while exercising on a treadmill.

response variable:

the number of calories burned

explanatory variables:

- ❑ duration of exercise (minutes)
- ❑ percentage of time spent at their target heart rate
- ❑ average speed (mph)
- ❑ age (years), gender, and body mass index (BMI)

Questions:

- ❑ What's the relationship between exercise duration and calories burned? Is it linear or curvilinear?
- ❑ How does effort (the percentage of time at the target heart rate, the average walking speed) factor in?
- ❑ How many calories can a 30-year-old man with a BMI of 28.7 expect to burn if he walks for 45 minutes at an average speed of 4 miles per hour and stays within his target heart rate 80 percent of the time?

Effective regression analysis is an interactive, holistic process with many steps, and it involves many skills.

Given the data, we repeat the following steps until we are satisfied with the model/results:

- ❑ propose a model
- ❑ fit the model
- ❑ check the model assumptions

Different regression models

Regression methods include the followings:

- ❑ simple/multiple linear
- ❑ polynomial
- ❑ logistic
- ❑ poisson
- ❑ time series

OLS-based regression

Here we focus on regression methods called ordinary least squares (OLS) regression, including simple/multiple linear regression and polynomial regression.

In OLS regression, a quantitative dependent variable is predicted from a weighted sum of predictor variables, where the weights are parameters estimated from the data.

Linear model

$$Y_i = \beta_o + \beta_1 X_{1i} + \dots + \beta_k X_{ki} + \varepsilon_i, \quad i = 1, \dots, n.$$

Model assumptions:

- Normality — For fixed values of the independent variables, the dependent variable is normally distributed.
- Independence — The Y_i values are independent of each other.
- Linearity
- Homoscedasticity — The variance of Y_i values doesn't vary with the levels of the independent variables.

OLS selects model parameters (intercept and slopes) that minimize the difference between actual response values and those predicted by the model. Specifically, model parameters are selected to minimize the sum of squared residuals.

$$\sum_{i=1}^n (Y_i - \hat{Y}_i)^2 = (Y_i - \hat{\beta}_o - \hat{\beta}_1 X_{1i} - \dots - \hat{\beta}_k X_{ki})^2 = \sum_{i=1}^n \varepsilon_i^2.$$

Fitting regression line with lm()

In R, the basic function for fitting a linear model is `lm()`. The format is

```
myfit <- lm(formula, data)
```

- ❑ formula: the model, ususally it is $y \sim x_1 + \dots + x_k$.
- ❑ data: the data frame for regression analysis, includes y, x_1, \dots, x_k .
- ❑ myfit: a list containing various information about the regression results.

Table 8.2 Symbols commonly used in R formulas

Symbol	Usage
~	Separates response variables on the left from the explanatory variables on the right. For example, a prediction of y from x , z , and w would be coded $y \sim x + z + w$.
+	Separates predictor variables.
:	Denotes an interaction between predictor variables. A prediction of y from x , z , and the interaction between x and z would be coded $y \sim x + z + x:z$.
*	A shortcut for denoting all possible interactions. The code $y \sim x * z * w$ expands to $y \sim x + z + w + x:z + x:w + z:w + x:z:w$.
^	Denotes interactions up to a specified degree. The code $y \sim (x + z + w)^2$ expands to $y \sim x + z + w + x:z + x:w + z:w$.
.	A place holder for all other variables in the data frame except the dependent variable. For example, if a data frame contained the variables x , y , z , and w , then the code $y \sim .$ would expand to $y \sim x + z + w$.

-	A minus sign removes a variable from the equation. For example, $y \sim (x + z + w)^2 - x:w$ expands to $y \sim x + z + w + x:z + z:w$.
-1	Suppresses the intercept. For example, the formula $y \sim x - 1$ fits a regression of y on x , and forces the line through the origin at $x=0$.
I()	Elements within the parentheses are interpreted arithmetically. For example, $y \sim x + (z + w)^2$ would expand to $y \sim x + z + w + z:w$. In contrast, the code $y \sim x + I((z + w)^2)$ would expand to $y \sim x + h$, where h is a new variable created by squaring the sum of z and w .
<i>function</i>	Mathematical functions can be used in formulas. For example, $\log(y) \sim x + z + w$ would predict $\log(y)$ from x , z , and w .

Table 8.3 Other functions that are useful when fitting linear models

Function	Action
<code>summary()</code>	Displays detailed results for the fitted model
<code>coefficients()</code>	Lists the model parameters (intercept and slopes) for the fitted model
<code>confint()</code>	Provides confidence intervals for the model parameters (95 percent by default)
<code>fitted()</code>	Lists the predicted values in a fitted model
<code>residuals()</code>	Lists the residual values in a fitted model
<code>anova()</code>	Generates an ANOVA table for a fitted model, or an ANOVA table comparing two or more fitted models
<code>vcov()</code>	Lists the covariance matrix for model parameters
<code>AIC()</code>	Prints Akaike's Information Criterion
<code>plot()</code>	Generates diagnostic plots for evaluating the fit of a model
<code>predict()</code>	Uses a fitted model to predict response values for a new dataset

A data example

The dataset women in the base installation provides the height and weight for a set of 15 women ages 30 to 39. We want to predict weight from height. Having an equation for predicting weight from height can help us to identify overweight or underweight individuals.

Example 1 (simple regression)

```
women
```

```
fit <- lm(weight ~ height, data=women)
```

```
summary(fit)
```

```
coefficients(fit)
```


```
confint(fit)
```

```
fitted(fit)
```

```
residuals(fit)
```

```
plot(women)
```

```
lines(women$height, fitted(fit))
```



Polynomial regression


The above plot suggests that you might be able to improve your prediction using a regression with a quadratic term (that is, X^2).

You can fit a quadratic equation using the statement

```
fit2 <- lm(weight ~ height + I(height^2), data=women)
```

Example 2 (polynomial regression)

```
fit2 <- lm(weight ~ height + I(height^2), data=women)
summary(fit2)
plot(women, xlab="Height (in inches)", ylab="Weight (in lbs)")
lines(women$height, fitted(fit2))
```



linear relationship

Linear versus nonlinear models

Note that this polynomial equation still fits under the rubric of linear regression. It's linear because the equation involves a weighted sum of predictor variables (height and height-squared in this case). Even a model such as

$$\hat{Y}_i = \hat{\beta}_0 \times \log X_1 + \hat{\beta}_2 \times \sin X_2$$

would be considered a linear model (linear in terms of the parameters) and fit with the formula

$$Y \sim \log(X1) + \sin(X2)$$

In contrast, here's an example of a truly nonlinear model:

$$Y_i = \hat{\beta}_0 + \hat{\beta}_1 e^{\frac{X}{\beta_2}}$$

Nonlinear models of this form can be fit with the `nls()` function.

Multiple linear regression

When there's more than one predictor variable, simple linear regression becomes multiple linear regression, and the analysis becomes more involved. We'll use the `state.x77` dataset in the base package for this example.

- ❑ response variable: a state's murder rate
- ❑ explanatory variables: population, illiteracy rate, average income, and frost levels (mean number of days below freezing).

Because the `lm()` function requires a data frame (and the `state.x77` dataset is a matrix), we need to subset the original data and transform it to a data frame.

Example 3 (data transformation)

```
state.x77
```

```
class(state.x77)
```

```
states <- as.data.frame(state.x77[,c("Murder", "Population",  
"Illiteracy", "Income", "Frost")])
```

```
states
```

```
class(states)
```

```
dim(states)
```

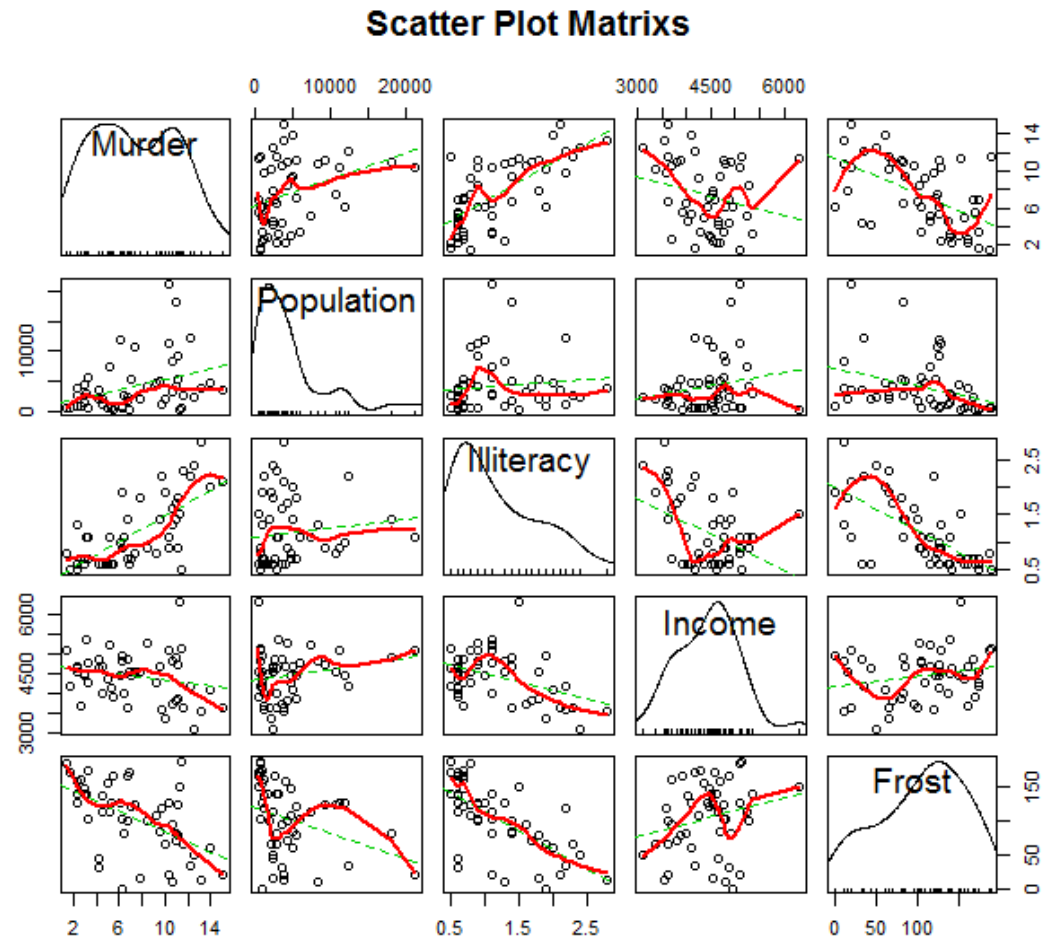
A good first step in multiple regression is to examine the relationships among the variables two at a time.

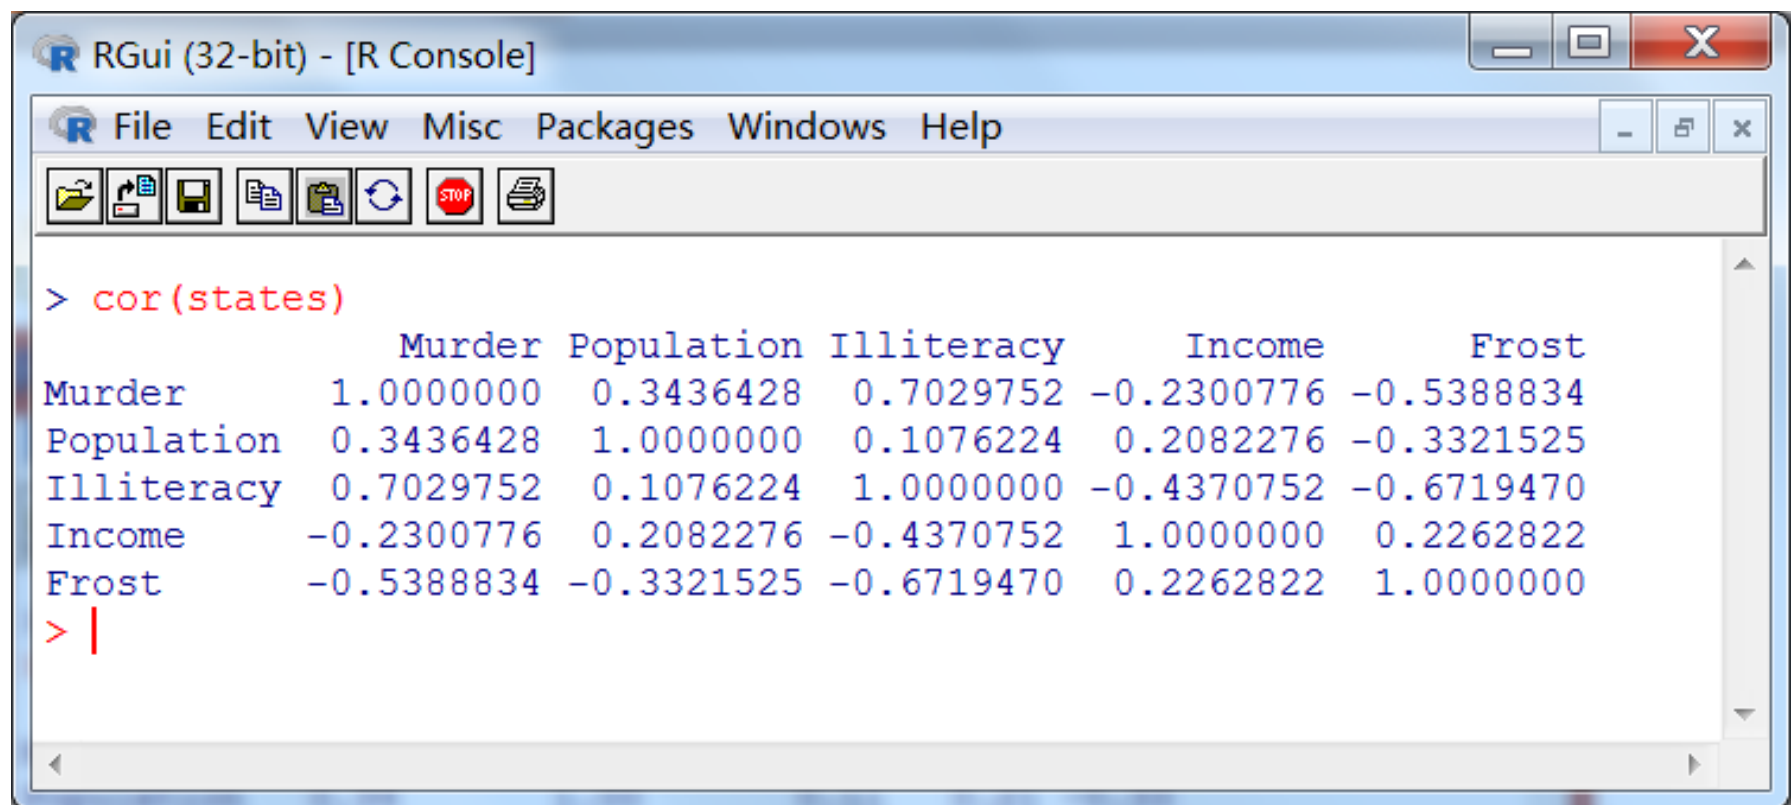
- scatterplot
- correlation matrix

Example 4 (multiple regression)

```
library(car)
scatterplotMatrix(states, spread=FALSE, lty=2, main="Scatter
Plot Matrixs")
options(digits=2)
cor(states)
```

Scatterplot





The screenshot shows the RGui (32-bit) - [R Console] window. The title bar includes the R logo and window controls. The menu bar contains File, Edit, View, Misc, Packages, Windows, and Help. The toolbar has icons for file operations and execution. The console displays the command `> cor(states)` and its output, a correlation matrix for five variables: Murder, Population, Illiteracy, Income, and Frost. The matrix is symmetric with 1.0000000 on the diagonal. The correlation between Murder and Frost is -0.5388834. The correlation between Population and Frost is -0.3321525. The correlation between Illiteracy and Frost is -0.6719470. The correlation between Income and Frost is 0.2262822. The correlation between Murder and Income is -0.2300776. The correlation between Population and Income is 0.2082276. The correlation between Illiteracy and Income is -0.4370752. The correlation between Murder and Population is 0.3436428. The correlation between Illiteracy and Population is 0.1076224.

```
> cor(states)
```

	Murder	Population	Illiteracy	Income	Frost
Murder	1.0000000	0.3436428	0.7029752	-0.2300776	-0.5388834
Population	0.3436428	1.0000000	0.1076224	0.2082276	-0.3321525
Illiteracy	0.7029752	0.1076224	1.0000000	-0.4370752	-0.6719470
Income	-0.2300776	0.2082276	-0.4370752	1.0000000	0.2262822
Frost	-0.5388834	-0.3321525	-0.6719470	0.2262822	1.0000000

```
> |
```

After we check that “murder rate” has medium or strong linear relationship with “population” and “ illiteracy”, we continue to do the linear regression.

Example 5 (linear regression)

```
fit <- lm(Murder ~ Population + Illiteracy + Income + Frost,  
data=states)
```

```
summary(fit)
```

```
coefficients(fit)
```

```
confint(fit)
```

```
residuals(fit)
```



```
> fit <- lm(Murder ~ Population + Illiteracy + Income + Frost, data=states)
> summary(fit)
```

Call:

```
lm(formula = Murder ~ Population + Illiteracy + Income + Frost,
    data = states)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.7960	-1.6495	-0.0811	1.4815	7.6210

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.235e+00	3.866e+00	0.319	0.7510
Population	2.237e-04	9.052e-05	2.471	0.0173 *
Illiteracy	4.143e+00	8.744e-01	4.738	2.19e-05 ***
Income	6.442e-05	6.837e-04	0.094	0.9253
Frost	5.813e-04	1.005e-02	0.058	0.9541

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.535 on 45 degrees of freedom

Multiple R-squared: 0.567, Adjusted R-squared: 0.5285

F-statistic: 14.73 on 4 and 45 DF, p-value: 9.133e-08



```
> coefficients(fit)
```

```
(Intercept)    Population    Illiteracy      Income      Frost
1.2345634112  0.0002236754  4.1428365903  0.0000644247  0.0005813055
```

```
> confint(fit)
```

```
                2.5 %      97.5 %
(Intercept) -6.552191e+00  9.0213182149
Population    4.136397e-05  0.0004059867
Illiteracy    2.381799e+00  5.9038743192
Income        -1.312611e-03  0.0014414600
Frost         -1.966781e-02  0.0208304170
```

```
> residuals(fit)
```

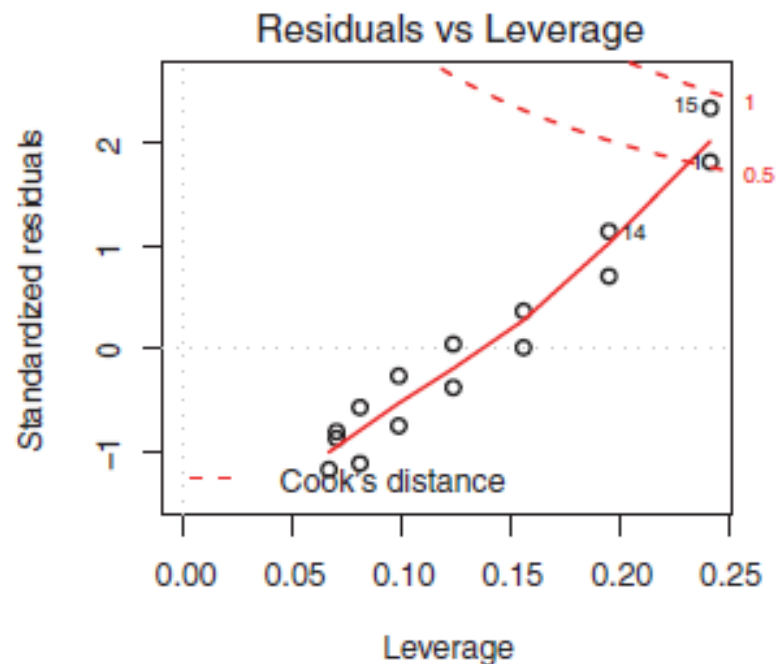



Alabama	Alaska	Arizona	Arkansas	California
4.11179210	3.27433977	-1.68700264	0.26668056	-0.57424792
Colorado	Connecticut	Delaware	Florida	Georgia
1.68594493	-3.81042204	0.73768277	1.91178879	2.97838044
Hawaii	Idaho	Illinois	Indiana	Iowa
-3.41984294	1.05927673	2.42954793	1.41893921	-2.02545720
Kansas	Kentucky	Louisiana	Maine	Maryland
-0.09731294	1.68494109	-0.72117551	-2.00277259	2.21479548
Massachusetts	Michigan	Minnesota	Mississippi	Missouri
-4.15834611	3.72023253	-2.69149081	0.57035176	3.34806321
Montana	Nebraska	Nevada	New Hampshire	New Jersey
0.74271628	-1.53684814	7.62104160	-1.39312273	-2.63613735
New Mexico	New York	North Carolina	North Dakota	Ohio

Regression diagnostics

`summary()` function outputs the model parameters and other statistics. But it does not tell us if the model we fit is appropriate.

Our confidence in inference of regression parameters depends on the degree to which the data have met the statistical assumptions of the OLS model.

```
> fit <- lm(weight ~ height, data=women)
> par(mfrow=c(2,2))
> plot(fit)
```



car package

Table 8.4 Useful functions for regression diagnostics (**car** package)

Function	Purpose
<code>qqPlot()</code>	Quantile comparisons plot
<code>durbinWatsonTest()</code>	Durbin–Watson test for autocorrelated errors
<code>crPlots()</code>	Component plus residual plots
<code>ncvTest()</code>	Score test for nonconstant error variance
<code>spreadLevelPlot()</code>	Spread-level plot
<code>outlierTest()</code>	Bonferroni outlier test
<code>avPlots()</code>	Added variable plots
<code>influencePlot()</code>	Regression influence plot
<code>scatterplot()</code>	Enhanced scatter plot
<code>scatterplotMatrix()</code>	Enhanced scatter plot matrix
<code>vif()</code>	Variance inflation factors


Normality

The `qqPlot()` function provides an accurate method of assessing the normality assumption.

Example 6 (checking normality)

```
library(car)
fit <- lm(Murder ~ Population + Illiteracy + Income + Frost, data=states)
qqPlot(fit, distribution="norm", labels=row.names(states),
simulate=TRUE, main="Q-Q Plot")
```

QQ plot



Independence of errors

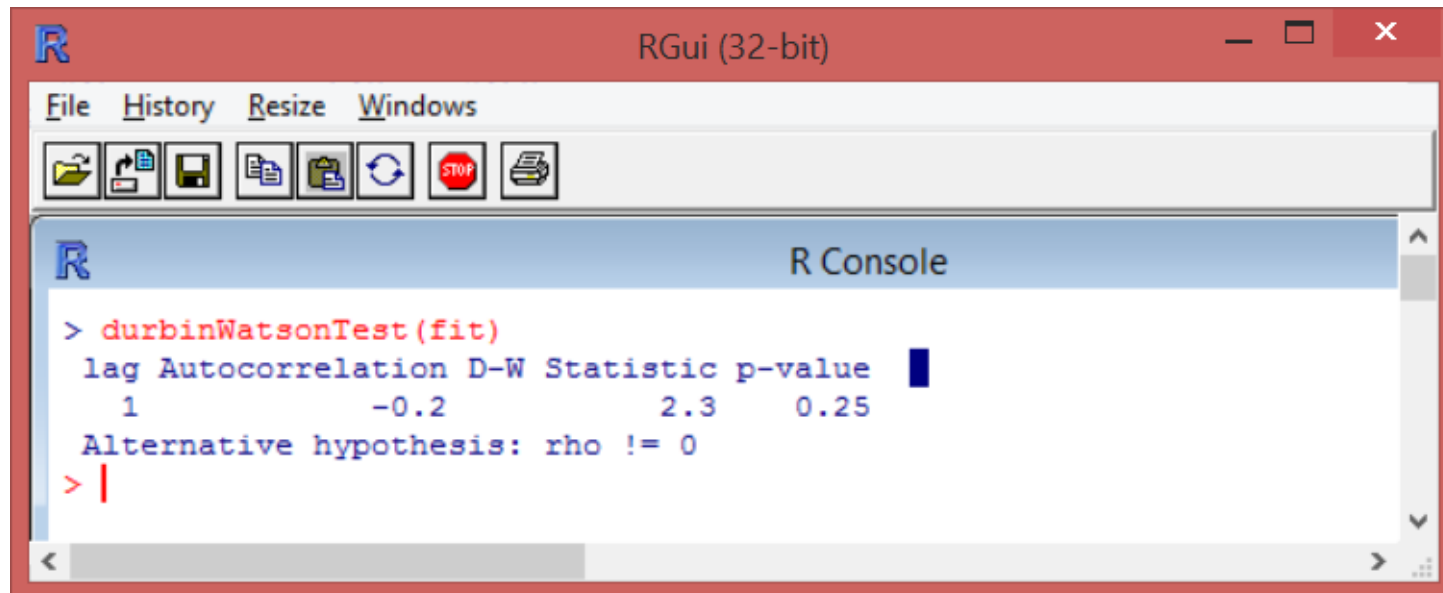
Time series data will often display autocorrelation. Observations collected closer in time will be more correlated with each other than with observations distant in time. We can apply the Durbin–Watson test to the multiple regression problem.

$$DW = \frac{(\varepsilon_2 - \varepsilon_1)^2 + \dots + (\varepsilon_n - \varepsilon_{n-1})^2}{\sum_{i=1}^n \varepsilon_i^2} \approx 2(1 - r)$$

Example 6 (checking independence)

`durbinWatsonTest(fit)`

Output:



The screenshot shows the RGui (32-bit) window. The title bar is red and contains the R logo, the text "RGui (32-bit)", and standard window controls. Below the title bar is a menu bar with "File", "History", "Resize", and "Windows". Under the menu bar is a toolbar with icons for file operations. The main area is the "R Console", which has a blue header bar with the R logo and the text "R Console". The console displays the output of the command `> durbinWatsonTest(fit)`. The output is a table with four columns: "lag", "Autocorrelation", "D-W", and "Statistic". The first row of data shows lag 1 with an autocorrelation of -0.2, a D-W statistic of 2.3, and a p-value of 0.25. Below the table, it says "Alternative hypothesis: rho != 0". The prompt `> |` is at the bottom of the console.

```
> durbinWatsonTest(fit)
lag Autocorrelation D-W Statistic p-value
  1          -0.2         2.3      0.25
Alternative hypothesis: rho != 0
> |
```

Homoscedasticity

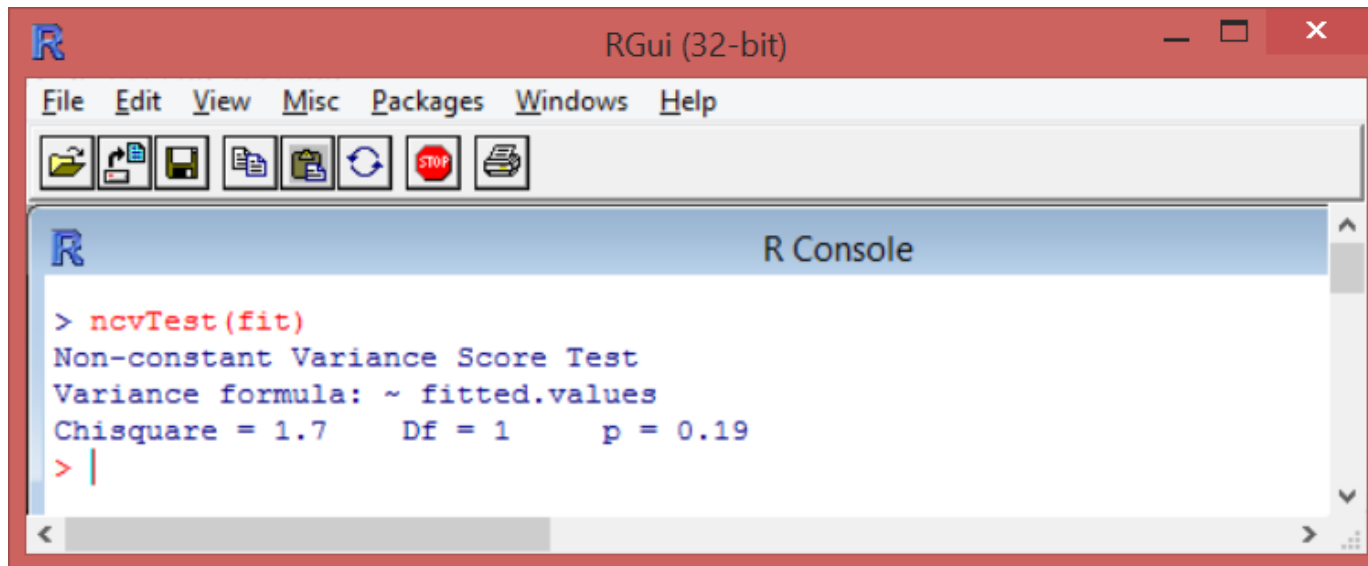
The car package also provides two useful functions for identifying non-constant error variance.

- ▣ `ncvTest()`
- ▣ `spreadLevelPlot()`

Example 7 (checking homoscedasticity)


`ncvTest(fit)`

`spreadLevelPlot(fit)`



```
RGui (32-bit)
File Edit View Misc Packages Windows Help
[Icons]
R Console
> ncvTest(fit)
Non-constant Variance Score Test
Variance formula: ~ fitted.values
Chisquare = 1.7    Df = 1    p = 0.19
> |
```

**Spread-Level Plot for
fit**



Summary

In this session, we have learned

- ❑ How to fit linear/polynomial/multiple regression model.
- ❑ Check the regression results by using `summary()` function.
- ❑ How to test the model assumptions: normality, independence and homoscedasticity.